

Install Production Kubernetes Cluster with Rancher RKE | ComputingForGeeks

Josphat Mutai

18-22 minutes

How can I use RKE to deploy Production ready Kubernetes Cluster?. Kubernetes has gained much traction and is now the standard orchestration layer for containerized workloads. If you want an open-source system for automating deployment of containerized applications without worrying about scaling and management, then Kubernetes is the right tool for you.

There are many standard ways of deploying a production-grade Kubernetes cluster. This include the use of tools such as [kops](#), [kubespray](#) or manually building a cluster with **kubeadm**. We have some of the guides you can use for reference.

[Deploy Production Ready Kubernetes Cluster with Ansible & Kubespray](#)

[Deploy Kubernetes Cluster on CentOS 7 / CentOS 8 With Ansible and Calico CNI](#)

This guide walks you through the simple steps for installation a production-grade Kubernetes cluster with RKE. We'll set up a 5-node cluster with Rancher Kubernetes Engine (RKE) and install the *Rancher chart* with the Helm package manager.

What is RKE?

Rancher Kubernetes Engine (RKE), is an extremely simple, lightning fast Kubernetes distribution that runs entirely within containers. Rancher is a container management platform built for organizations that deploy containers in production. Rancher makes it easy to run Kubernetes everywhere, meet IT requirements, and empower DevOps teams.

Prepare Workstation machine

A number of CLI Tools are required on your Workstation where deployment will be done. This can also be a virtual machine that is able to access cluster nodes.

1. kubectl:

```
--- Linux ---
curl -LO https://storage.googleapis.com/kubernetes-release/release/`curl -s https://storage.googleapis.com/kubernetes-release/release/stable.txt`/bin/linux/amd64/kubectl
chmod +x ./kubectl
sudo mv ./kubectl /usr/local/bin/kubectl
kubectl version --client

--- macOS ---
curl -LO "https://storage.googleapis.com/kubernetes-release/release/\$\(curl -s https://storage.googleapis.com/kubernetes-release/release/stable.txt\)/bin/darwin/amd64/kubectl"
chmod +x ./kubectl
sudo mv ./kubectl /usr/local/bin/kubectl
kubectl version --client
```

2. rke

```
--- Linux ---
```

```
curl -s https://api.github.com/repos/rancher/rke/releases/latest | grep
download_url | grep amd64 | cut -d '"' -f 4 | wget -qi -
chmod +x rke_linux-amd64
sudo mv rke_linux-amd64 /usr/local/bin/rke
rke --version
```

```
--- macOS ---
curl -s https://api.github.com/repos/rancher/rke/releases/latest | grep
download_url | grep darwin-amd64 | cut -d '"' -f 4 | wget -qi -
chmod +x rke_darwin-amd64
sudo mv rke_darwin-amd64 /usr/local/bin/rke
rke --version
```

3. helm

```
--- Helm 3 ---
curl -fsSL -o get_helm.sh https://raw.githubusercontent.com/helm/helm/master/scripts/get-helm-3
chmod 700 get_helm.sh
./get_helm.sh
```

Install Kubernetes with RKE

I'll be working on 5 nodes:

- **3 Master Nodes** – etcd and control plane (3 for HA)
- **2 Worker nodes** – Scale to meet your Workloads demand

These are the specifications for my setup.

- Master Nodes – **8GB** of RAM and **4** vcpus
- Worker Machines – **16GB** of RAM and **8** vpcus

RKE Supported operating systems

RKE runs on almost any Linux OS with Docker installed. Rancher has been tested and is supported with:

- Red Hat Enterprise Linux
- Oracle Enterprise Linux
- CentOS Linux
- Ubuntu
- RancherOS

Step 1: Update your Linux System

The first step is to update your Linux machines which will be used to build the cluster.

```
--- CentOS ---
$ sudo yum -y update
$ sudo reboot

--- Ubuntu / Debian ---
$ sudo apt-get update
$ sudo apt-get upgrade
$ sudo reboot
```

Step 2: Create rke user

If using Red Hat Enterprise Linux, Oracle Enterprise Linux or CentOS, you cannot use the root user as [SSH user](#) due to [Bugzilla 1527565](#). For this reason, we'll create a user account called **rke** for deployment purpose.

Using Ansible Playbook:

```
---
- name: Create rke user with passwordless sudo
  hosts: rke-hosts
  remote_user: root
  tasks:
    - name: Add RKE admin user
      user:
        name: rke
        shell: /bin/bash

    - name: Create sudo file
      file:
        path: /etc/sudoers.d/rke
        state: touch

    - name: Give rke user passwordless sudo
      lineinfile:
        path: /etc/sudoers.d/rke
        state: present
        line: 'rke ALL=(ALL:ALL) NOPASSWD: ALL'

    - name: Set authorized key taken from file
      authorized_key:
        user: rke
        state: present
        key: "{{ lookup('file', '~/.ssh/id_rsa.pub') }}"
```

Create user manually on all hosts

Login to each of your cluster nodes and create *rke* user.

```
sudo useradd rke
sudo passwd rke
```

Enable passwordless sudo for the user:

```
$ sudo vim /etc/sudoers.d/rke
rke ALL=(ALL:ALL) NOPASSWD: ALL
```

Copy your ssh public key to the user's *~/.ssh/authorized_keys* file.

```
for i in rke-master-01 rke-master-02 rke-master-03 rke-worker-01 rke-
worker-02; do
  ssh-copy-id rke@$i
done
```

Confirm you can login from your workstation:

```
$ ssh rke@rke-master-01
Warning: Permanently added 'rke-master-01,x.x.x.x' (ECDSA) to the list
of known hosts.
[rke@rke-master-01 ~]$ sudo su - # No password prompt
Last login: Mon Jan 27 21:28:53 CET 2020 from y.y.y.y on pts/0
[root@rke-master-01 ~]# exit
[rke@rke-master-01 ~]$ exit
logout
Connection to rke-master-01 closed.
```

Step 3: Enable required Kernel modules:

Using Ansible:

Create a playbook with below contents and run it against your RKE servers inventory.

```
---
```

```

- name: Load RKE kernel modules
  hosts: rke-hosts
  remote_user: root
  vars:
    kernel_modules:
      - br_netfilter
      - ip6_udp_tunnel
      - ip_set
      - ip_set_hash_ip
      - ip_set_hash_net
      - iptable_filter
      - iptable_nat
      - iptable_mangle
      - iptable_raw
      - nf_conntrack_netlink
      - nf_conntrack
      - nf_conntrack_ipv4
      - nf_defrag_ipv4
      - nf_nat
      - nf_nat_ipv4
      - nf_nat_masquerade_ipv4
      - nfnetlink
      - udp_tunnel
      - veth
      - vxlan
      - x_tables
      - xt_addrtype
      - xt_conntrack
      - xt_comment
      - xt_mark
      - xt_multiport
      - xt_nat
      - xt_recent
      - xt_set
      - xt_statistic
      - xt_tcpudp

  tasks:
    - name: Load kernel modules for RKE
      modprobe:
        name: "{{ item }}"
        state: present
      with_items: "{{ kernel_modules }}"

```

The manual way

Login to each host and enable Kernel modules required to run Kubernetes.

```

for module in br_netfilter ip6_udp_tunnel ip_set ip_set_hash_ip
ip_set_hash_net iptable_filter iptable_nat iptable_mangle iptable_raw
nf_conntrack_netlink nf_conntrack nf_conntrack_ipv4 nf_defrag_ipv4
nf_nat nf_nat_ipv4 nf_nat_masquerade_ipv4 nfnetlink udp_tunnel veth
vxlan x_tables xt_addrtype xt_conntrack xt_comment xt_mark xt_multiport
xt_nat xt_recent xt_set xt_statistic xt_tcpudp;
do
    if ! lsmod | grep -q $module; then
        echo "module $module is not present";
    fi;
done

```

Step 4: Disable swap and Modify sysctl entries

The recommendation of Kubernetes is to disable swap and add some sysctl values.

With Ansible:

```

---
- name: Disable swap and load kernel modules
  hosts: rke-hosts
  remote_user: root
  tasks:
    - name: Disable SWAP since kubernetes can't work with swap enabled
      (1/2)
      shell: |
        swapoff -a

    - name: Disable SWAP in fstab since kubernetes can't work with swap
      enabled (2/2)
      replace:
        path: /etc/fstab
        regexp: '^([^#].*?\sswap\s+.*)$'
        replace: '# \1'
    - name: Modify sysctl entries
      sysctl:
        name: '{{ item.key }}'
        value: '{{ item.value }}'
        sysctl_set: yes
        state: present
        reload: yes
      with_items:
        - {key: net.bridge.bridge-nf-call-ip6tables, value: 1}
        - {key: net.bridge.bridge-nf-call-iptables, value: 1}
        - {key: net.ipv4.ip_forward, value: 1}

```

Manually

Swap:

```
$ sudo vim /etc/fstab
# Add comment to swap line
```

```
$ sudo swapoff -a
```

Sysctl:

```
$ sudo tee -a /etc/sysctl.d/99-kubernetes.conf <<EOF
net.bridge.bridge-nf-call-iptables = 1
net.ipv4.ip_forward = 1
net.bridge.bridge-nf-call-ip6tables = 1
EOF
```

```
$ sysctl --system
```

Confirm is disabled:

```
$ free -h
```

	total	used	free	shared	buff/cache
available					
Mem:	7.6G	180M	6.8G	8.5M	633M
7.2G					
Swap:	0B	0B	0B		

Step 5: Install Supported version of Docker

Each Kubernetes version supports different Docker versions. The Kubernetes release notes contain the [current list](#) of validated Docker versions.

As of this article, supported docker versions are:

Docker Version	Install Script
19.03	<code>curl https://releases.rancher.com/install-docker/19.03.sh sh</code>

Docker Version	Install Script
18.09	<code>curl https://releases.rancher.com/install-docker/18.09.sh sh</code>
18.06	<code>curl https://releases.rancher.com/install-docker/18.06.sh sh</code>
17.03	<code>curl https://releases.rancher.com/install-docker/17.03.sh sh</code>

You can either follow the [Docker installation](#) instructions or use one of Rancher's [install scripts](#) to install Docker. I'll install the latest supported version:

```
curl https://releases.rancher.com/install-docker/19.03.sh | sudo bash -
```

Start and enable docker service:

```
sudo systemctl enable --now docker
```

Confirm that a Kubernetes supported version of Docker is installed on your machine:

```
$ sudo docker version --format '{{.Server.Version}}'
18.09.2
```

Add **rke** user to docker group.

```
$ sudo usermod -aG docker rke
$ id rke
uid=1000(rke) gid=1000(rke) groups=1000(rke),994(docker)
```

Step 6: Open Ports on firewall

- **For a single-node installation**, you only need to open the ports required to enable Rancher to communicate with downstream user clusters.
- **For a high-availability installation**, the same ports need to be opened, as well as additional ports required to set up the Kubernetes cluster that Rancher is installed on.

Check all ports used in the [requirements page](#)

Firewalld TCP ports:

```
for i in 22 80 443 179 5473 6443 8472 2376 8472 2379-2380 9099 10250
10251 10252 10254 30000-32767; do
    sudo firewall-cmd --add-port=${i}/tcp --permanent
done
sudo firewall-cmd --reload
```

Firewalld UDP ports:

```
for i in 8285 8472 4789 30000-32767; do
    sudo firewall-cmd --add-port=${i}/udp --permanent
done
```

Step 6: Allow SSH TCP Forwarding

You need to enable your SSH server system-wide TCP forwarding.

Open ssh configuration file located at **/etc/ssh/sshd_config**:

```
$ sudo vi /etc/ssh/sshd_config
AllowTcpForwarding yes
```

Restart ssh service after making the change.

```
--- CentOS ---
$ sudo systemctl restart sshd
```

```
--- Ubuntu ---
$ sudo systemctl restart ssh
```

Step 7: Generate RKE cluster configuration file.

RKE uses a cluster configuration file, referred to as `cluster.yml` to determine what nodes will be in the cluster and how to deploy Kubernetes.

There are [many configuration options](#) that can be set in the `cluster.yml`. This file can be created from [minimal example](#) templates or generated with the **rke config** command.

Run `rke config` command to create a new `cluster.yml` in your current directory.

```
rke config --name cluster.yml
```

This command will prompt you for all the information needed to build a cluster.

If you want to create an empty template `cluster.yml` file instead, specify the `--empty` flag.

```
rke config --empty --name cluster.yml
```

This is how my cluster configuration file looks like – **Don't copy paste, just use it as reference** to create your own configuration.

```
# https://rancher.com/docs/rke/latest/en/config-options/
nodes:
- address: 10.10.1.10
  internal_address:
  hostname_override: rke-master-01
  role: [controlplane, etcd]
  user: rke
- address: 10.10.1.11
  internal_address:
  hostname_override: rke-master-02
  role: [controlplane, etcd]
  user: rke
- address: 10.10.1.12
  internal_address:
  hostname_override: rke-master-03
  role: [controlplane, etcd]
  user: rke
- address: 10.10.1.13
  internal_address:
  hostname_override: rke-worker-01
  role: [worker]
  user: rke
- address: 10.10.1.114
  internal_address:
  hostname_override: rke-worker-02
  role: [worker]
  user: rke

# using a local ssh agent
# Using SSH private key with a passphrase - eval `ssh-agent -s` && ssh-add
ssh_agent_auth: true

# SSH key that access all hosts in your cluster
ssh_key_path: ~/.ssh/id_rsa

# By default, the name of your cluster will be local
# Set different Cluster name
cluster_name: rke

# Fail for Docker version not supported by Kubernetes
ignore_docker_version: false

# prefix_path: /opt/custom_path

# Set kubernetes version to install: https://rancher.com/docs/rke/latest/en/upgrades/#listing-supported-kubernetes-versions
# Check with -> rke config --list-version --all
```

```

kubernetes_version:
# Etcd snapshots
services:
  etcd:
    backup_config:
      interval_hours: 12
      retention: 6
    snapshot: true
    creation: 6h
    retention: 24h

kube-api:
# IP range for any services created on Kubernetes
# This must match the service_cluster_ip_range in kube-controller
service_cluster_ip_range: 10.43.0.0/16
# Expose a different port range for NodePort services
service_node_port_range: 30000-32767
pod_security_policy: false

kube-controller:
# CIDR pool used to assign IP addresses to pods in the cluster
cluster_cidr: 10.42.0.0/16
# IP range for any services created on Kubernetes
# This must match the service_cluster_ip_range in kube-api
service_cluster_ip_range: 10.43.0.0/16

kubelet:
# Base domain for the cluster
cluster_domain: cluster.local
# IP address for the DNS service endpoint
cluster_dns_server: 10.43.0.10
# Fail if swap is on
fail_swap_on: false
# Set max pods to 150 instead of default 110
extra_args:
  max-pods: 150

# Configure network plug-ins
# KE provides the following network plug-ins that are deployed as add-
ons: flannel, calico, weave, and canal
# After you launch the cluster, you cannot change your network provider.
# Setting the network plug-in
network:
  plugin: canal
  options:
    canal_flannel_backend_type: vxlan

# Specify DNS provider (coredns or kube-dns)
dns:
  provider: coredns

# Currently, only authentication strategy supported is x509.
# You can optionally create additional SANs (hostnames or IPs) to
# add to the API server PKI certificate.
# This is useful if you want to use a load balancer for the
# control plane servers.
authentication:
  strategy: x509
  sans:
    - "k8s.computingforgeeks.com"

# Set Authorization mechanism
authorization:
  # Use `mode: none` to disable authorization

```



```

mode: rbac

# Currently only nginx ingress provider is supported.
# To disable ingress controller, set `provider: none`
# `node_selector` controls ingress placement and is optional
ingress:
  provider: nginx
  options:
    use-forwarded-headers: "true"

```

In my configuration, the master nodes only have **etcd** and **controlplane** roles. But they can be used to schedule pods by adding **worker** role.

```
role: [controlplane, etcd, worker]
```

Step 7: Deploy Kubernetes Cluster with RKE

Once you've created the **cluster.yml** file, you can deploy your cluster with a simple command.

```
rke up
```

This command assumes the *cluster.yml* file is in the same directory as where you are running the command. If using a different filename, specify it like below.

```
$ rke up --config ./rancher_cluster.yml
```

Using SSH private key with a passphrase – `eval ssh-agent -s && ssh-add`

Ensure the setup doesn't show any failure in its output:

```

.....
INFO[0181] [sync] Syncing nodes Labels and Taints
INFO[0182] [sync] Successfully synced nodes Labels and Taints
INFO[0182] [network] Setting up network plugin: canal
INFO[0182] [addons] Saving ConfigMap for addon rke-network-plugin to
Kubernetes
INFO[0183] [addons] Successfully saved ConfigMap for addon rke-network-
plugin to Kubernetes
INFO[0183] [addons] Executing deploy job rke-network-plugin
INFO[0189] [addons] Setting up coredns
INFO[0189] [addons] Saving ConfigMap for addon rke-coredns-addon to
Kubernetes
INFO[0189] [addons] Successfully saved ConfigMap for addon rke-coredns-
addon to Kubernetes
INFO[0189] [addons] Executing deploy job rke-coredns-addon
INFO[0195] [addons] CoreDNS deployed successfully..
INFO[0195] [dns] DNS provider coredns deployed successfully
INFO[0195] [addons] Setting up Metrics Server
INFO[0195] [addons] Saving ConfigMap for addon rke-metrics-addon to
Kubernetes
INFO[0196] [addons] Successfully saved ConfigMap for addon rke-metrics-
addon to Kubernetes
INFO[0196] [addons] Executing deploy job rke-metrics-addon
INFO[0202] [addons] Metrics Server deployed successfully
INFO[0202] [ingress] Setting up nginx ingress controller
INFO[0202] [addons] Saving ConfigMap for addon rke-ingress-controller to
Kubernetes
INFO[0202] [addons] Successfully saved ConfigMap for addon rke-ingress-
controller to Kubernetes
INFO[0202] [addons] Executing deploy job rke-ingress-controller
INFO[0208] [ingress] ingress controller nginx deployed successfully
INFO[0208] [addons] Setting up user addons
INFO[0208] [addons] no user addons defined
INFO[0208] Finished building Kubernetes cluster successfully

```

Step 8: Accessing your Kubernetes cluster

As part of the Kubernetes creation process, a kubeconfig file has been created and written at kube_config_cluster.yml.

Set KUBECONFIG variable to the file generated.

```
export KUBECONFIG=./kube_config_cluster.yml
```

Check list of nodes in the cluster.

```
$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
rke-master-01	Ready	controlplane,etcd	16m	v1.17.0
rke-master-02	Ready	controlplane,etcd	16m	v1.17.0
rke-master-03	Ready	controlplane,etcd	16m	v1.17.0
rke-worker-01	Ready	worker	6m33s	v1.17.0
rke-worker-02	Ready	worker	16m	v1.17.0

You can copy this file to **\$HOME/.kube/config** if you don't have any other kubernetes cluster.

Step 9: Install Kubernetes Dashboard

If you'll prefer deploying containerized applications on Kubernetes from a dashboard, use our guide below.

[How To Install Kubernetes Dashboard with NodePort](#)

Learning Courses:

In our next guide, we'll cover installation of Rancher – An **open-source multi-cluster orchestration platform** that lets you easily manage and secure your enterprise Kubernetes.

We really appreciate you supporting our efforts by buying us Coffee:

Coming up with fresh, high quality content takes time. Sometimes working late at night building labs and then doing the writing. We appreciate if you consider supporting our efforts with a cup of coffee to keep us awake and always deliver.



No contribution is small. We are grateful for any amount you support us with. Thank you!